

# Atomic Workshop

# Lazy initialization

- We saw earlier in the course how to perform the following lazy initialization with thread-safety by using double-checked locking with `std::call_once`

```
some_type *ptr{ nullptr };    // Variable to be lazily initialized

void process() {
    if (!ptr)                  // First time variable has been used
        ptr = new some_type;  // Initialize it
    ptr->do_it();              // Use it
}
```

- In this exercise, make the lazy initialization thread-safe by using an atomic type
- Write a simple program to exercise your code

# Lock-free queue

- Consider the following code to implement a simple queue without any internal or external locks
- The queue is only accessed by two threads
  - A Producer thread inserts elements into the queue
  - A Consumer thread removes elements from the queue
- The code is carefully designed so that the Consumer and Producer threads never work on adjacent elements
  - This ensures that the two threads always work on different parts of the queue

# Lock-free queue

- The queue has two iterators, iHead and iTail
  - iHead points to the element before the oldest element
  - iTail points to the element before the newest (most recently added)
  - When the Producer thread adds an element, it increments iTail
  - When the Consumer thread removes an element, it increments iHead
- Only the Producer thread modifies the queue
  - As well as inserting elements, the Producer queue is responsible for erasing elements that the Consumer thread has removed
- The Producer thread never erases the iHead element
  - This maintains separation between the threads

# Lock-free queue class

```
template <typename T>
struct LockFreeQueue {
    private:
        std::list<T> list;
        typename std::list<T>::iterator iHead, iTail;

    public:
        LockFreeQueue() {
            list.push_back(T());           // Add dummy separator
            iHead = list.begin();
            iTail = list.end();
        }
}
```

# Producer task member function

```
void Produce(const T& t) {  
    list.push_back(t);           // Add the new item  
    iTail = list.end();          // Publish it  
    list.erase(list.begin(), iHead); // Trim unused nodes  
}
```

# Consumer task member function

```
bool Consume(T& t) {  
    auto iNext = iHead;  
    ++iNext;  
    if (iNext != iTail) {  
        iHead = iNext;  
        t = *iHead;  
        return true;  
    }  
    return false;  
}  
}; // End of class definition
```

// If queue is not empty  
// Publish that we took an item  
// Copy it back to the caller  
// And report success

// Else report queue was empty

# Exercise

- Why is the "dummy separator necessary?
- Add a member function to print out all the elements
- Add a main function which calls Produce and Consume in separate threads
- Write a loop which runs the Produce and Consume threads and, once they have completed, calls Print
- Increase the number of iterations until you observe a race condition
- Explain why the race condition occurs
- Can the race condition be avoided by using atomic variables?